

STA-380 Exercises

Khyati Jariwala, Aishwarya Rajeev, Juhi Patel, Varun Kausika

2022-08-10

1. Probability practice

Part A

Given information:

Two categories of users:

1. Truthful clicker (TC)
2. Random clicker (RC)

Information on probabilities:

- $P(RC) = 0.3$
- $P(Yes|RC) = 0.5$
- $P(No|RC) = 0.5$
- $P(TC) = 0.7$
- $P(Yes|TC) = x$
- $P(No|TC) = 1 - x$
- $P(Yes) = 0.65$
- $P(No) = 0.35$

Using the Rule of Total Probability,

$$P(Yes) = P(Yes, TC) + P(Yes, RC) = P(TC) * P(Yes|TC) + P(RC) * P(Yes|RC) \quad (1)$$

$$P(Yes) = 0.7x + 0.3 * 0.5 = 0.7x + 0.15 = 0.65$$

Solving for x, we get,

$$x = P(Yes|TC) = 0.714$$

Part B

We are being asked $P(Diseased|Positive)$

Given information:

- $P(Positive|Diseased) = 0.993$
- $P(Negative|NotDiseased) = 0.9999$
- $P(Diseased) = 0.000025$

According to Bayes Rule and Rule of Total Probability,

$$P(\text{Diseased}|\text{Positive}) = \frac{P(\text{Positive}|\text{Diseased}) * P(\text{Diseased})}{P(\text{Positive})} \quad (2)$$

and,

$$P(\text{Positive}) = P(\text{Positive}|\text{Diseased}) * P(\text{Diseased}) + P(\text{Positive}|\text{Not Diseased}) * P(\text{Not Diseased}) \quad (3)$$

Therefore,

$$P(\text{Positive}) = 0.993 * 0.000025 + 0.0001 * 0.999975 = 0.000125$$

Substituting in (2) we get,

$$P(\text{Diseased}|\text{Positive}) = \frac{0.993 * 0.000025}{0.000125} = 0.1986$$

2. Wrangling the Billboard Top 100

Part A

First, we load in the data and perform a group by on the performer and the song, with an agg function of count for the week

Table 1: Billboards

performer	song	count
'N Sync	(God Must Have Spent) A Little More Time On You	22
'N Sync	Bye Bye Bye	23
'N Sync	Gone	24
'N Sync	I Drive Myself Crazy	12
'N Sync	I Want You Back	24
'N Sync	It's Gonna Be Me	25
'N Sync	Pop	15
'N Sync	Tearin' Up My Heart	1
'N Sync	This I Promise You	26
'N Sync & Gloria Estefan	Music Of My Heart	20

Finally, we sort the dataframe in descending order of counts and find the top 10 and give our table a caption:

Table 2: Top 10 most popular songs

performer	song	count
Imagine Dragons	Radioactive	87
AWOLNATION	Sail	79
Jason Mraz	I'm Yours	76
The Weeknd	Blinding Lights	76
LeAnn Rimes	How Do I Live	69

performer	song	count
LMFAO Featuring Lauren Bennett & GoonRock	Party Rock Anthem	68
OneRepublic	Counting Stars	68
Adele	Rolling In The Deep	65
Jewel	Foolish Games/You Were Meant For Me	65
Carrie Underwood	Before He Cheats	64

Part B

First we group by year and order by ascending year. Then, we remove the years 1958 and 2021 from the rows and order just to make sure. Finally, we proceed to plot the columns.

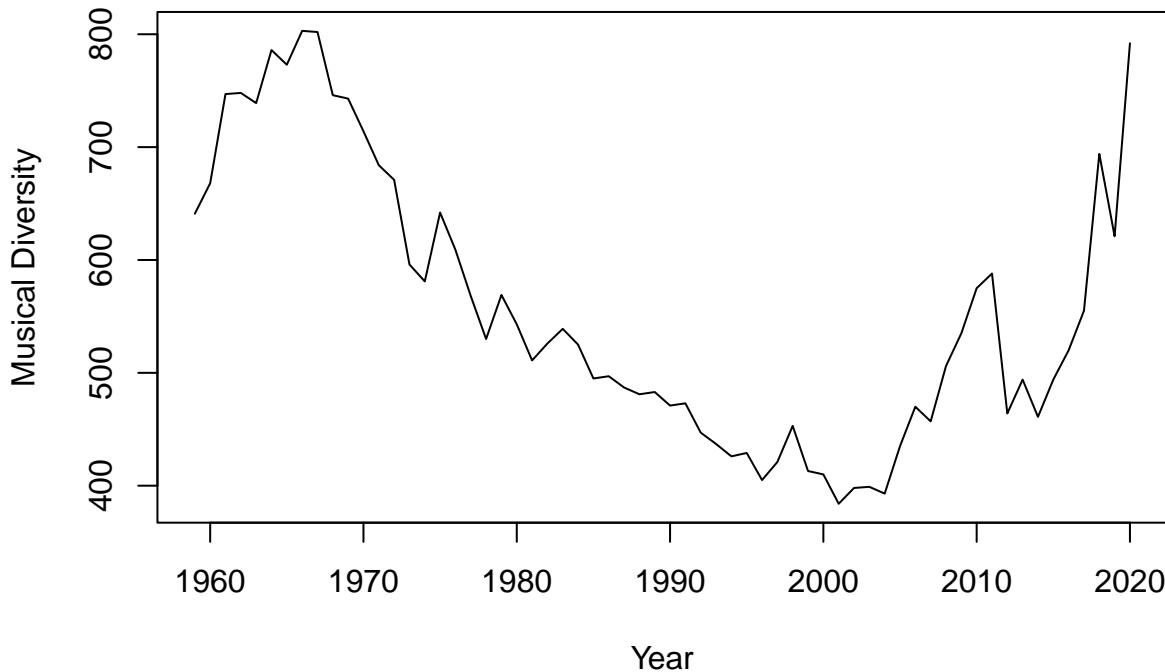


Figure 1: The plot shows peaks in diversity in 1965 and 2020, along with extreme lows in 2000

Part C

First, we filter the dataframe from part A to include only those songs with weeks at least 10. Then, We do a group by on the artists. Finally, we can select those artists with a hit-count of at least 30. ??

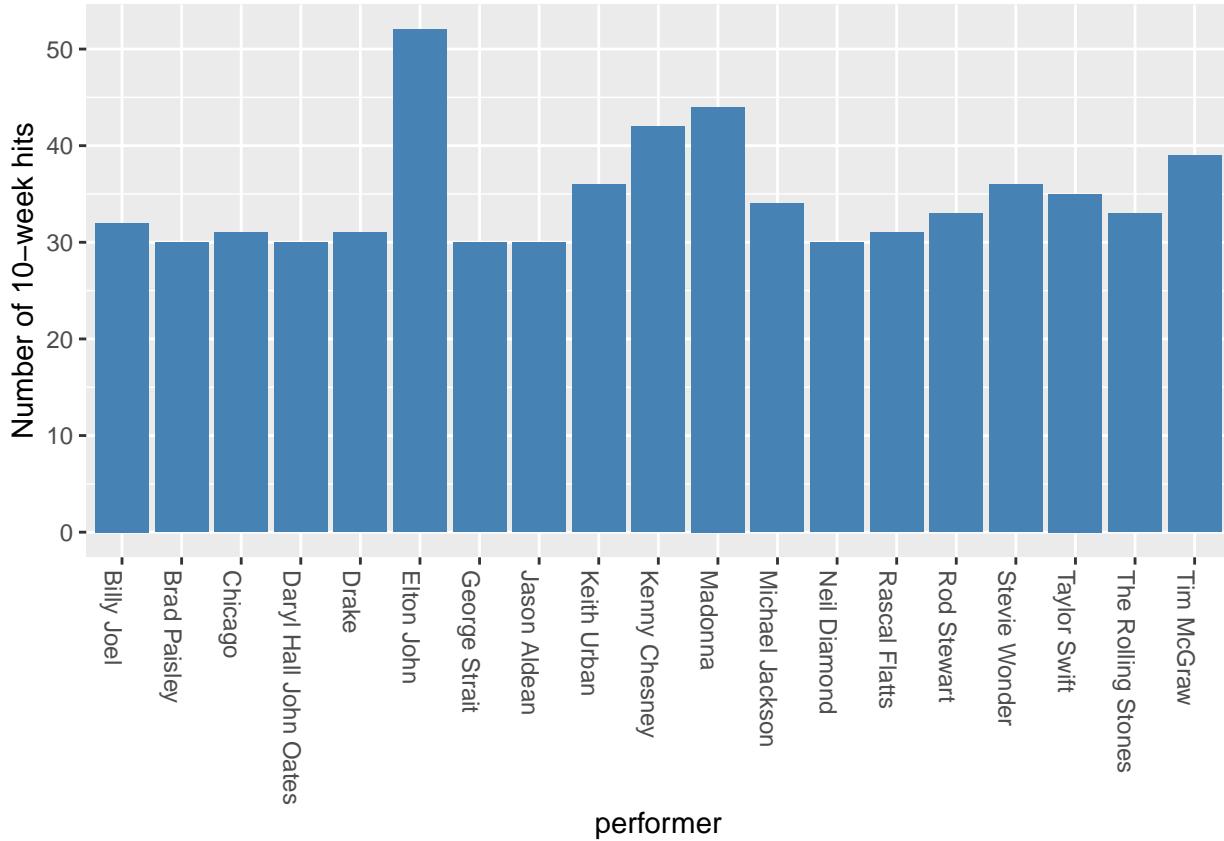


Figure 2: Elton John has more hits than others by quite a large margin

3. Visual Storytelling Part 1: green buildings

Outliers:

We are interested in finding the potential economic gains the owner could make by constructing a 15-story green building in the neighbourhood of East Cesar Chavez.

First, we remove the outliers from the dataset as mentioned by the analyst.




Figure 3: It seems like occupancy below 40 percent can be removed in contrast to the 10 percent that was suggested

Observations:

1. Removing outliers in occupancy below 40% would only remove 456 rows from our dataset, so we proceed with this step.
2. We should not remove outliers from rent, as that would shrink our dataset greatly, and also since most of the green buildings have higher rents, this would remove them entirely.
3. We should not remove outliers from size, because size and rent are positively correlated and this might remove a significant number of green buildings as well.

Correlation Matrix:

Now, we plot the feature correlation matrix in our dataset.

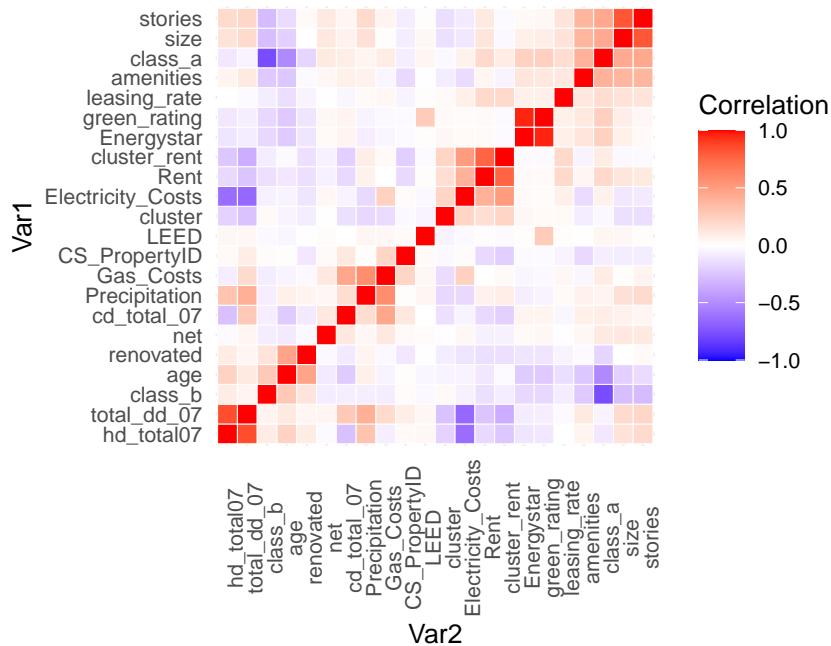


Figure 4: As we can see, rent is highly dependent on clusters

Observations:

1. Rent and Cluster has a high positive correlation, because some neighbourhoods are more expensive than others. So, instead of taking the median rent throughout our dataset, we can just use the cluster rent variable.
2. Rent and age has a mildly negative correlation, as one would expect.
3. Rent is not correlated with green rating, LEED or Energy Star.
4. Class A buildings are positively correlated with rent whereas Class B buildings are negatively correlated.
5. Leasing rate is positively correlated with rent.
6. Size and stories(can be considered correlated features) are both positively correlated to rent.

Scatter plots:

Now, we draw some scatter plots to understand the relation between correlated variables.

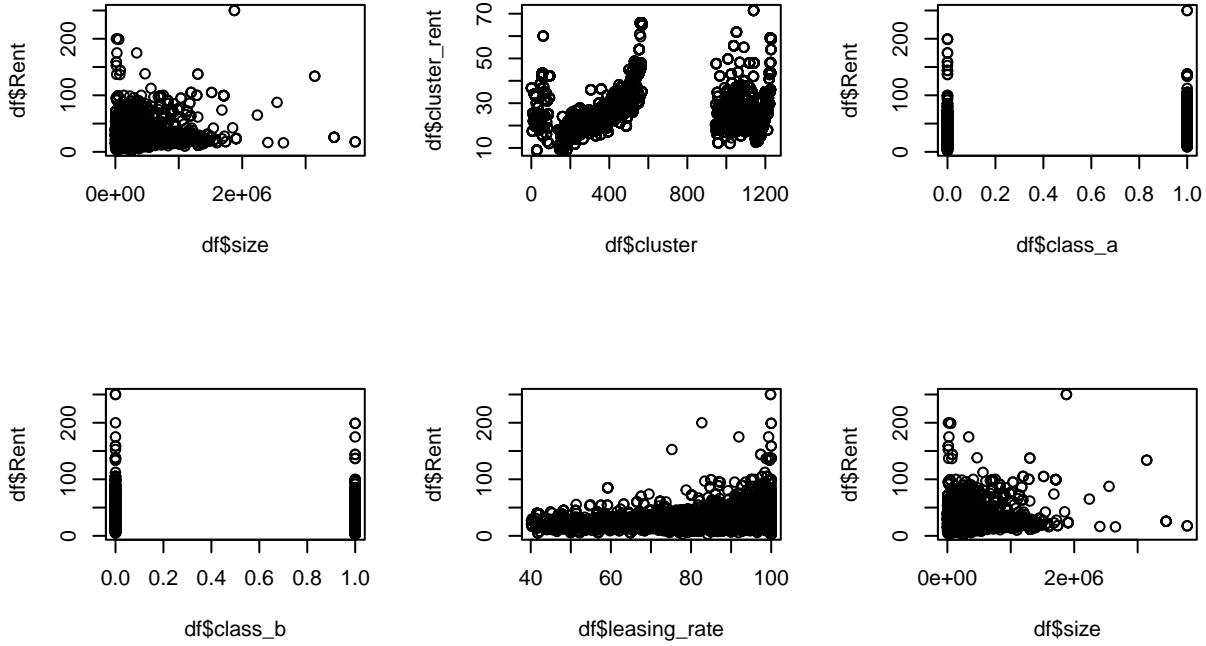


Figure 5: We can see clear patterns in the cluster rent v cluster. Some clusters have a significantly higher average rent.

Observations

1. Rent in Class A buildings occur more in positive class. This could be because of the positive correlation of class A buildings with green rating, size and stories.
2. Rent in Class B buildings occur more in negative class. The correlation of Class B with the variables is exactly opposite to that of Class A.
3. This implies that the builder **MUST** aim to construct a Class A building in order to be able to charge more rent.
4. Furthermore, the builder might want to move the building location to a cluster that falls in higher rent bracket.

Density plots:

Now, we look at the green buildings in our dataset and compare them to all the other buildings

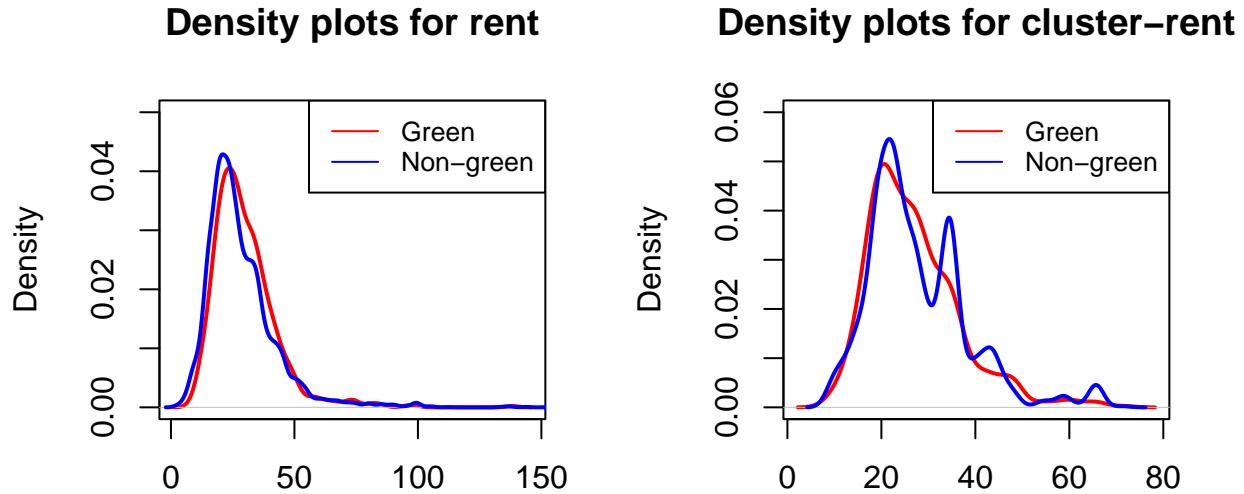


Figure 6: We can see that both green buildings and non-green buildings are not normally distributed. They are shifted to the left.

Observations:

1. We can conclude that this data is skewed toward the lower rent buildings (has long tails).
2. When comparing green and non-green buildings, the green buildings seem to peak at a slightly higher rent than the non green rents. Comparing the median values of both of these will be appropriate for the graph on the left.
3. However, for the cluster rent graph, we see that the peak of the non-green occurs at a slightly higher rent value than the peak for the green buildings. Therefore, if we use cluster rent as our primary variable to determine income (as we said we would before), we would expect in fact a lower rent value for the greens.
4. The cluster rent density function is not normal for either the greens or the non greens. For the non-greens, it looks more like a sum of many different bell shaped curves. This suggests that each cluster within the data has a separate normal distribution.

Economic Impact of Green Houses:

Economic impact could include a variety of factors in this case, namely:

1. **Sources of income depend on the rent being charged. Which inturn are caused by:**
 - (a) Size of the house being rented (250,000 sqft in our case).
 - (b) The neighbourhood of the house. A groupby statement is in order to find out the prices in East Cesar Chavez. On top of this, the premium rent for being green, as shown in the density plot, is pretty much no-existent. So controlling for the clusters, we can say that green houses don't do much better than non-green houses in this dataset.
 - (c) Age of the building. One would expect that the older the building is the lower the rent is. This is shown in the correlation plot.

- (d) Appreciation of value of houses in the neighbourhood. This is something that cannot be measured by the given features. A thorough analysis of previous time series data of prices is recommended.
- (e) Occupancy of the building (relevant variable here is lease.rate).
- (f) Whether amenities are available or not.
- (g) Whether it was renovated or not.
- (h) Whether it is Class A or Class B.

2. Sources of expenditure depend on:

- (a) Initial capex (100,000 in our case).
- (b) Premium needed to be spent on constructing a green building (5%).
- (c) Maintenance, repairs and other charges.
- (d) **NOTE: Water, Gas and Electricity charges are assumed to be paid by the tenant, hence are not included.**

Where the Analyst went wrong:

The analyst found the impact on income in a very linear way. He did not take into account the following:

1. That the green buildings do not have any premium in rent after accounting for the cluster in which the buildings are.
2. That 40% of the occupancies are outliers, since they have a very high mean and low interquartile range.
3. Time value of money: he has calculated the simple Payback Period of the project. The more accurate estimator of success would have been the Net Present Value.
4. A cross sectional analysis across other buildings should have been done to compare Payback Periods instead of assuming 8 years is a decent time to recuperate costs
5. The worst case of occupancy is not 90%. It is much lower. As seen in the dataset many buildings have occupancies even lower than 40%.
6. That the cost of the buildings in the area may reduce/increase over time.
7. Accounting for other variables which may inhibit the ability to charge high rent (eg. Class A vs B, others listed above).

Confounding variables:

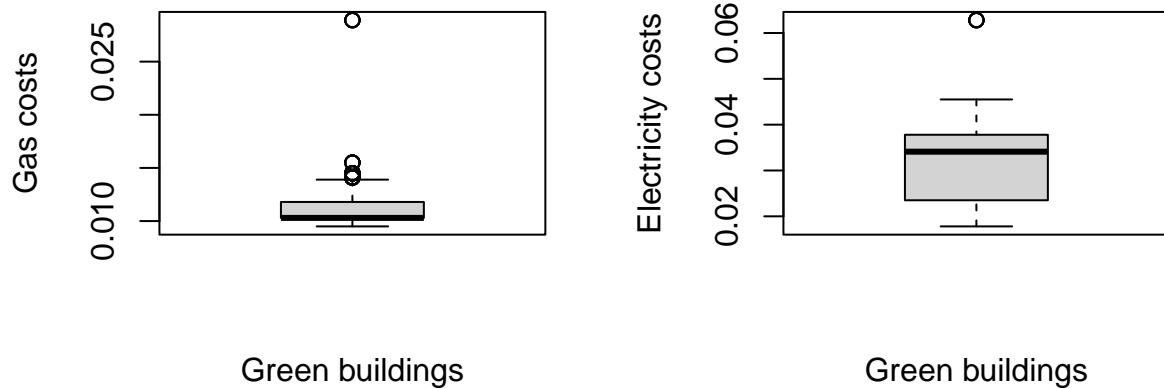


Figure 7: Some green houses show a high electricity and gas cost which goes against our intuition of what is green. The outliers can be adjusted for by removing them

4. Visual Storytelling part 2: Capital Metro data:

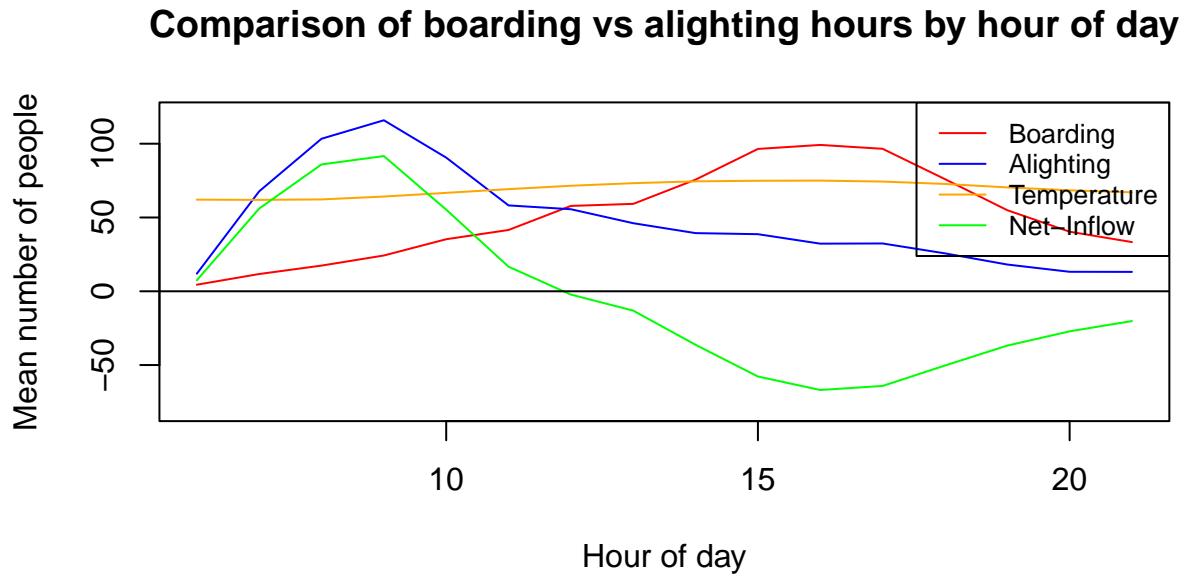


Figure 8: In the campus vicinity, alighting is most common in the morning (when classes start) and boarding is most common in the evening (when classes end). Furthermore, temperature doesn't seem to have an affect on the boardings and alightings. Here, net-inflow is calculated as alightings minus boardings. Net inflow is 0 around noon (break between classes).

Comparison of boarding vs alighting means by month and hour of day

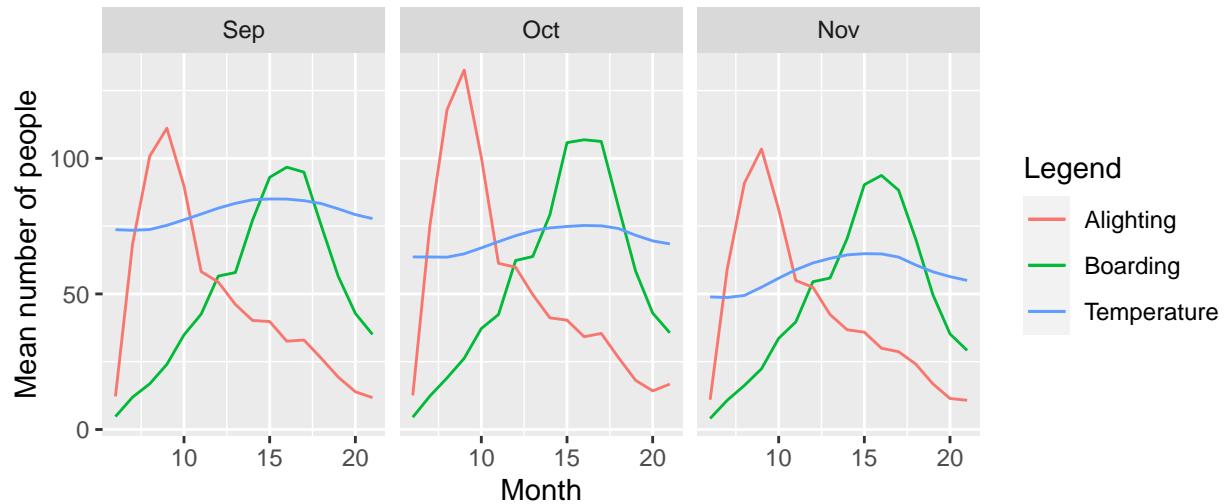


Figure 9: Similar patterns are seen across months, temperatures are highest in September and gradually come down

Comparison of boarding vs alighting means by day of week




Figure 10: As we can see, the number of people boarding and alighting near campus on weekends are much lower.

Comparison of boarding vs alighting means by month and day of week

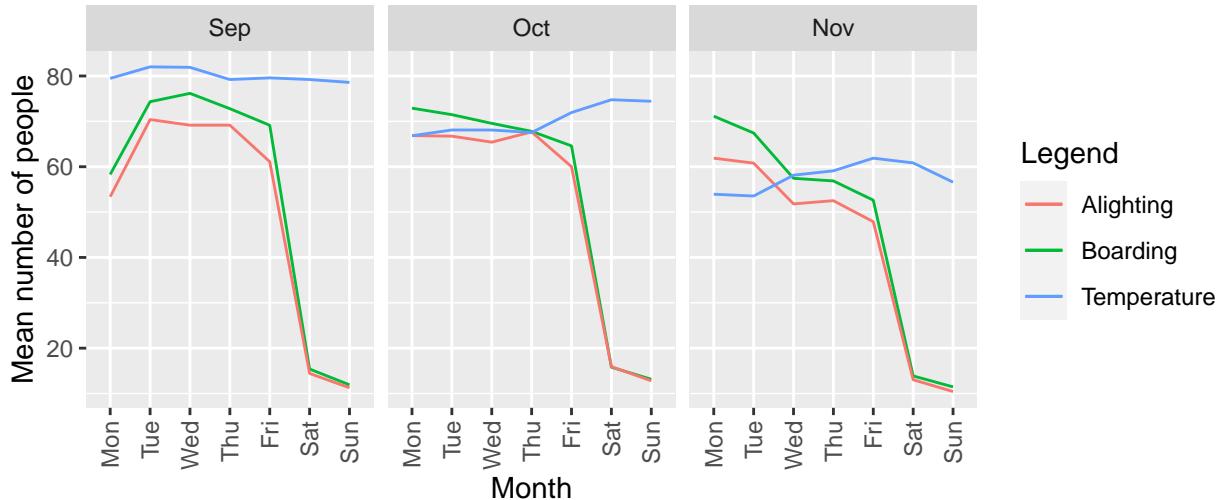


Figure 11: Similar patterns are seen across months, however significantly less people seem to attend classes on Monday in September.

Comparison of boarding vs alighting means by month




Figure 12: There seem to be more students coming in October compared to November and September. This is because of Thanksgiving holidays in November and since college starts in August, students might not immediately be using the buses on a regular basis in September.

5. Portfolio modeling

Introduction

We chose the below ETFs to provide portfolio diversification and a range of risk.

- Invesco (QQQ): One of the biggest, exclusively non-financial stock, and heavily tech-heavy trusts.

- SPY: One of the safest and largest ETF.
- iShares Russell 1000 Growth ETF (IWF): One of the most popular US large-cap growth ETFs with a long track record. It is an aggressive growing ETF.
- SCO: A low performing stock.

In total, we have selected 3 ETFs - “QQQ”, “SPY”, “IWF”. We looked at data from ETFs for five years commencing on January 1, 2017

Sample Data for QQQ

Table 3: QQQ Sample Data

QQQ.Open	QQQ.High	QQQ.Low	QQQ.Close	QQQ.Volume	QQQ.Adjusted
114.6697	115.3619	114.3043	114.9293	22307600	114.9293
115.0542	115.7657	115.0446	115.5542	19749100	115.5542
115.4677	116.2849	115.4677	116.2080	20644300	116.2080
116.3330	117.5347	116.0349	117.2271	24074300	117.2271
117.3232	117.8232	117.2463	117.6117	18909200	117.6117
117.6597	118.2270	117.3905	117.8712	16176600	117.8713
117.9193	118.1885	117.3136	118.1885	20686000	118.1885
117.7655	118.0828	116.9098	118.0058	19622500	118.0058
118.0635	118.6115	118.0635	118.4096	16930800	118.4097
118.1020	118.2847	117.7078	118.0539	14538500	118.0539

Combine all the returns in a matrix

Table 4: All Returns

	C1C1.QQQa	C1C1.SPYa	C1C1.IWFa	C1C1.SCOa
2017-01-04	0.0054375	0.0059492	0.0077585	-0.0220811
2017-01-05	0.0056577	-0.0007945	0.0015961	-0.0204144
2017-01-06	0.0087697	0.0035778	0.0064680	0.0026839
2017-01-09	0.0032806	-0.0033009	-0.0000932	0.0606204
2017-01-10	0.0022071	0.0000000	0.0003726	0.0429038
2017-01-11	0.0026917	0.0028261	0.0007449	-0.0530961
2017-01-12	-0.0015456	-0.0025099	-0.0010234	0.9404691
2017-01-13	0.0034219	0.0022955	0.0029803	0.0220018
2017-01-17	-0.0030042	-0.0035235	-0.0015787	-0.0012128
2017-01-18	0.0020360	0.0022099	0.0019532	0.0428051

Compute the returns from the closing prices

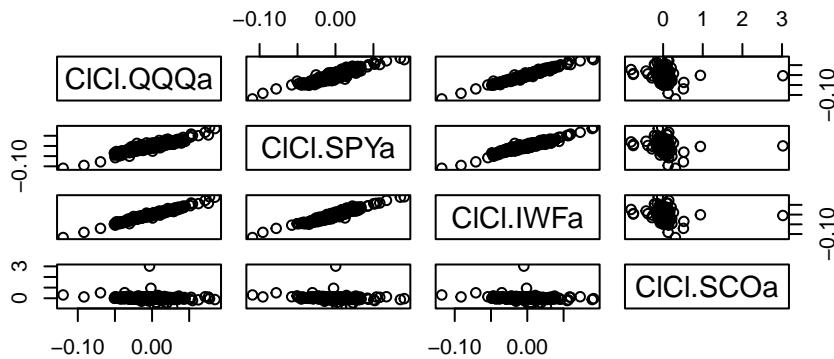


Figure 13: We can see a strong linear correlation here in 3 of the stocks QQQ, SPY and IWF. SCO seems to not have a strong correlation with the other stocks

Volatility of the ETFs

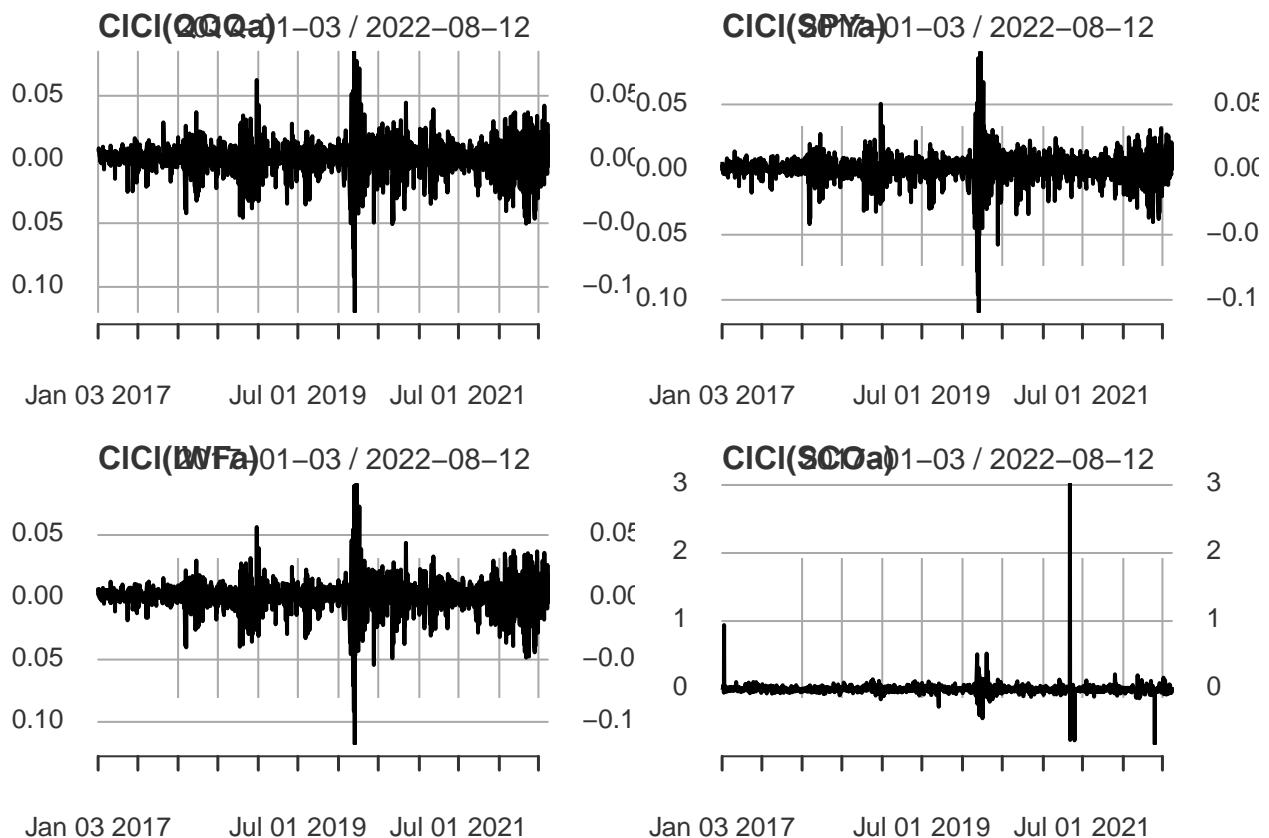


Figure 14: Volatility of the ETFs across the 5 year period.

Sample a random return from the empirical joint distribution

We simulated a random day's return.

```
##          C1C1.QQQa  C1C1.SPYa  C1C1.IWFa  C1C1.SCOa
## 2020-11-16  0.00780263  0.01248255  0.00525364 -0.0415355
```

SIMULATION 1: LOW RISK PORTFOLIO

- INITIAL INVESTMENT IS \$ 100000.
- Average return of investment after 20 days = \$ 101190.
- 5% value at Risk for safe portfolio = \$ 12139.7.

Now simulate many different possible futures, repeating the above block thousands of times

```
##          [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## result.1  100025.03  99877.23  100193.04  100374.26  99625.71  96487.93  97413.44
## result.2   97645.18  97019.54  96239.56  97457.63  96709.73  97313.76  97977.01
## result.3   99517.79  99492.94  101159.87  100681.74  101423.88  103183.56  103096.60
## result.4  101442.00  100935.60  103448.92  102039.49  103123.97  104118.65  104762.17
## result.5   99474.78  100467.10  101100.43  100714.50  101333.30  100132.13  98823.24
## result.6  100263.80  100834.35  100448.19  98329.78  98085.05  97114.86  98392.72
## result.7  100328.18  96501.36  96639.18  95979.78  97169.25  96641.63  96628.56
## result.8   98975.37  100279.24  100088.42  99671.95  99929.20  99933.75  99888.97
## result.9   99905.17  100956.47  100581.08  101405.10  103603.35  102064.87  100580.00
## result.10  99016.41  98472.57  97838.74  98517.02  103810.87  103636.66  104349.01
##          [,8]      [,9]      [,10]     [,11]     [,12]     [,13]     [,14]
## result.1  100497.09  99921.46  99940.86  99408.28  98408.91  98566.87  102087.43
## result.2   98229.31  97115.23  96891.40  96291.05  93609.72  93050.87  93353.84
## result.3  105078.23  104508.88  104390.85  105181.15  105231.31  105707.43  105131.59
## result.4  103812.60  104813.87  105877.52  105126.08  104543.94  105175.52  104733.71
## result.5  101026.30  101897.69  101160.75  100678.54  100731.01  102682.47  102651.92
## result.6   98707.96  99895.90  99181.88  97716.62  97625.80  97955.81  98290.08
## result.7   96875.78  96959.73  97005.34  96783.23  96795.33  97669.40  95895.67
## result.8   99888.89  102162.06  101998.81  101915.25  102177.09  102091.57  104594.90
## result.9  100204.84  99186.45  99399.80  100087.00  99920.35  101661.35  101132.93
## result.10  102939.02  102159.35  100771.05  100440.34  99674.24  101909.53  99559.33
##          [,15]     [,16]     [,17]     [,18]     [,19]     [,20]
## result.1   92692.72  92000.06  91914.92  90919.68  92442.49  93350.74
## result.2   93651.97  94333.45  93612.61  95906.43  98141.48  97279.79
## result.3  104886.14  96130.76  95229.86  95265.98  95383.47  96780.00
## result.4  105039.36  107250.54  107578.46  109943.38  109773.35  111440.80
## result.5  103268.40  103318.04  103244.54  102805.14  101819.84  100588.24
## result.6  101722.34  101197.23  100690.52  100525.19  99684.61  99044.00
## result.7   95383.83  95273.85  94962.64  95355.06  95114.06  93967.50
## result.8  104077.90  103740.78  104200.30  103394.30  102548.63  102295.37
## result.9  101608.46  101363.34  101078.19  101893.09  101444.90  98940.62
## result.10  100922.43  99985.79  100372.86  100986.85  101472.71  100511.51
```

Histogram of sim1[, n_days]

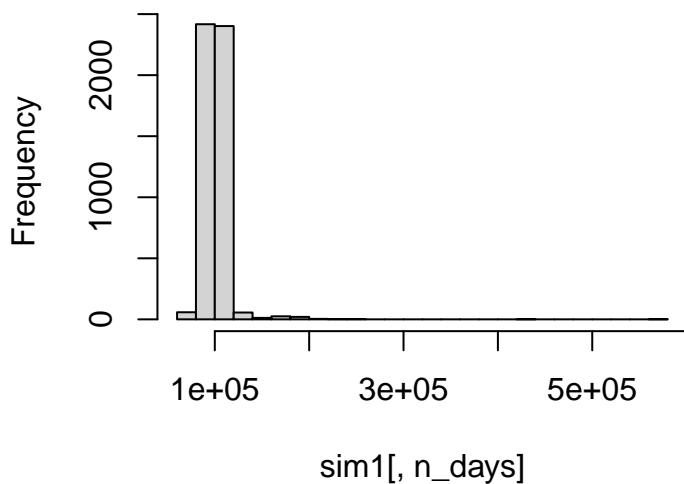


Figure 15: Histogram and density plots of simulation 1. We see spikes at 100000

density.default(x = sim1[, n_days])

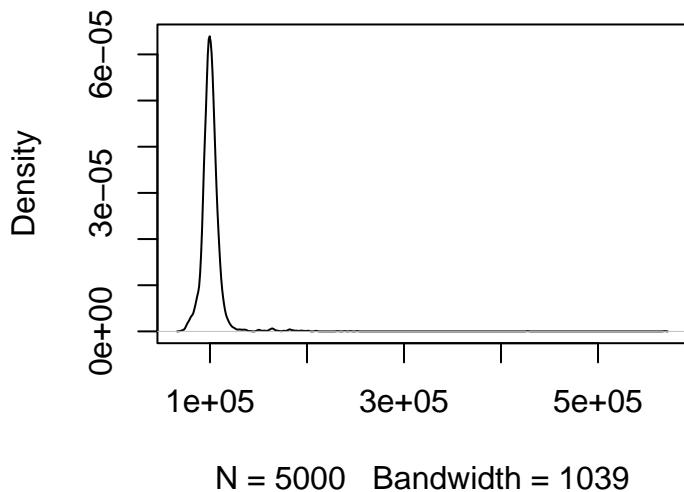


Figure 16: Histogram and density plots of simulation 1. We see spikes at 100000

Profit and Loss

- Average return of investment after 20 days = \$101099.5
- Average profit/loss after 20 days \$1099.513

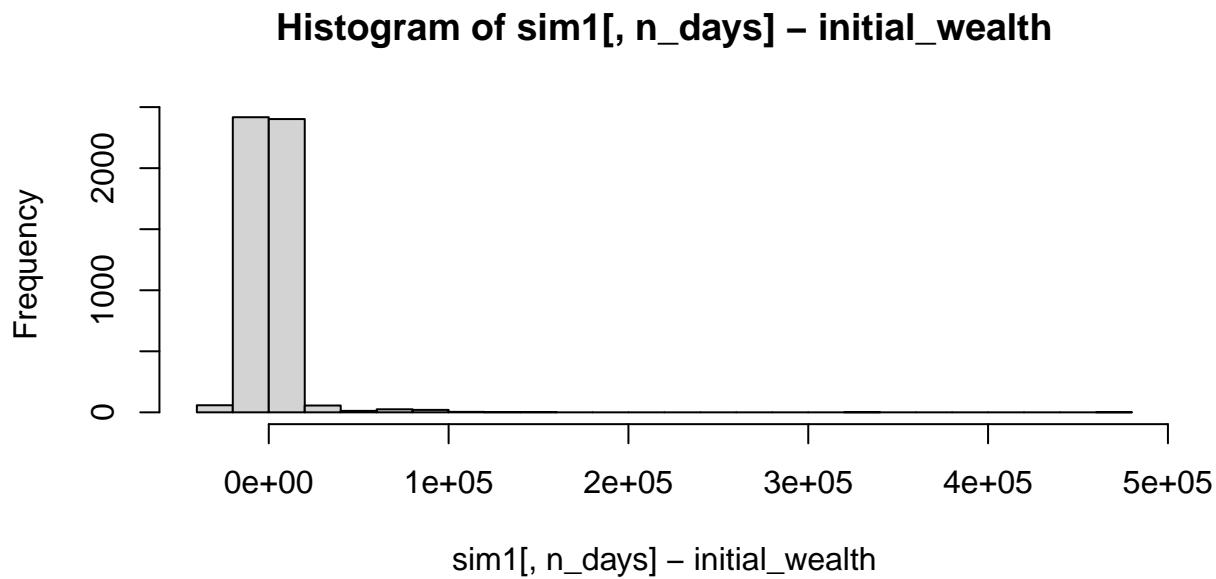


Figure 17: Histogram of returns of simulation 1.

5% Value at Risk for the first simulation = -12293.44

SIMULATION 2: HIGH RISK PORTFOLIO

- Average return of investment after 20 days = \$101576.2
- 5% Value at Risk for safe portfolio = \$7849.9327

```
##          [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## result.1 99907.76 100563.85 99311.09 99145.26 98258.92 97870.60 98096.91
## result.2 96421.68 95709.17 96480.34 97416.92 97205.12 97304.13 96559.88
## result.3 98847.61 98297.00 98897.86 99083.52 97876.25 97580.36 97554.75
## result.4 99555.86 99599.53 99222.58 100224.08 100049.35 100158.74 101383.49
## result.5 99291.32 97720.32 97720.23 96587.07 95845.40 96294.77 96396.97
## result.6 94178.60 94695.30 93407.15 93004.70 91907.62 93968.69 96135.22
## result.7 99823.12 101757.13 103057.75 103200.95 103786.53 104699.39 107407.01
## result.8 100026.06 100959.34 101280.08 101432.75 101672.86 103063.27 103388.54
## result.9 100325.59 100201.62 96635.85 96335.01 96193.59 98466.79 97886.42
## result.10 100079.63 100026.30 101068.65 101388.47 100124.58 101676.39 101367.76
##          [,8]      [,9]      [,10]     [,11]     [,12]     [,13]     [,14]
## result.1 97997.72 99506.52 99447.35 99147.45 99513.31 99233.34 98532.11
## result.2 96465.46 97009.71 97688.41 97726.21 98408.27 98850.31 98790.54
## result.3 98417.63 97376.23 97912.89 98587.93 99073.43 99380.47 99024.88
## result.4 100542.43 101598.71 100930.16 100635.26 100127.31 100339.57 102228.89
## result.5 97325.09 98023.56 98587.17 100522.08 100993.53 99316.36 99897.42
## result.6 96179.49 99046.87 99574.51 98005.82 97628.42 97487.29 97221.83
## result.7 106637.05 108026.19 108553.35 110403.42 110042.80 111043.26 109897.60
## result.8 104223.51 104399.94 104441.77 104545.24 103631.42 105103.41 106268.47
## result.9 98203.62 97103.66 96590.26 96045.98 94795.48 93839.58 93785.12
## result.10 100254.75 99451.33 100918.87 102087.72 101846.09 99467.33 100171.27
```

```

##          [,15]      [,16]      [,17]      [,18]      [,19]      [,20]
## result.1 98440.52 98742.37 98265.94 95761.65 96660.49 97118.14
## result.2 98513.06 97978.35 97369.96 97633.73 97007.64 97944.39
## result.3 100111.40 99759.35 99343.79 98541.59 98933.74 99881.66
## result.4 103643.75 106359.61 108125.64 109049.77 109310.45 109202.30
## result.5 100465.25 101331.88 98242.46 96647.14 96621.52 98259.03
## result.6 96937.26 103138.52 103176.94 103231.04 104727.83 102311.15
## result.7 109781.52 110642.35 108996.90 110233.33 110079.54 111536.82
## result.8 105889.01 106032.97 103113.24 103446.98 102708.44 103268.51
## result.9 93603.61 94289.91 93446.53 93208.80 93421.20 94320.68
## result.10 101232.21 100801.23 100884.21 101013.67 100608.26 101997.31

```

Histogram of sim2[, n_days]

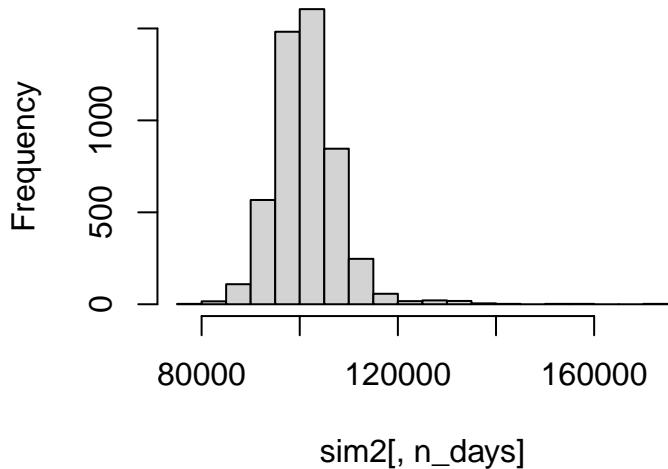


Figure 18: Histogram and density plots of simulation 2. We see spikes at 100000

density.default(x = sim2[, n_days])

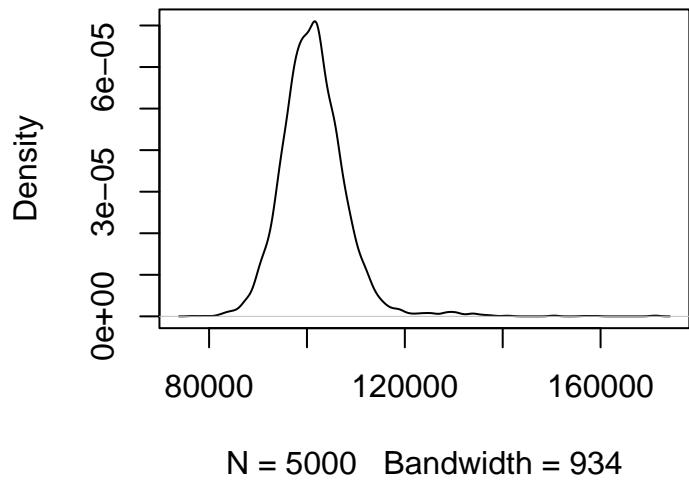


Figure 19: Histogram and density plots of simulation 2. We see spikes at 100000

Profit and Loss

Histogram of sim2[, n_days] – initial_wealth

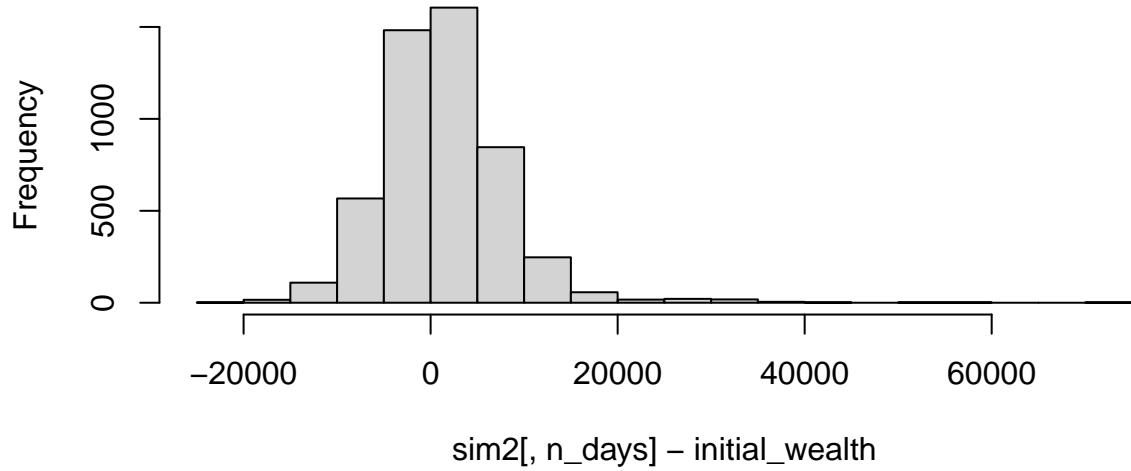


Figure 20: Histogram of returns of simulation 2.

5% Value at Risk for the second simulation = -8281.054

Summary

For the safe portfolio, we are observing lower return of investment and lower 5% VaR.

As the portfolio risk increased, we are able to witness the increase in returns and an increase in VaR value as expected.

References: <https://www.bankrate.com/investing/best-etfs/> <https://etfdb.com/compare/lowest-ytd-returns/>

Clustering and PCA

Sample Data

```
##   fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 1      7.4          0.70     0.00       1.9      0.076
## 2      7.8          0.88     0.00       2.6      0.098
## 3      7.8          0.76     0.04       2.3      0.092
## 4     11.2          0.28     0.56       1.9      0.075
## 5      7.4          0.70     0.00       1.9      0.076
## 6      7.4          0.66     0.00       1.8      0.075
##   free.sulfur.dioxide total.sulfur.dioxide density    pH sulphates alcohol
## 1              11            34 0.9978 3.51      0.56     9.4
## 2              25            67 0.9968 3.20      0.68     9.8
## 3              15            54 0.9970 3.26      0.65     9.8
## 4              17            60 0.9980 3.16      0.58     9.8
## 5              11            34 0.9978 3.51      0.56     9.4
## 6              13            40 0.9978 3.51      0.56     9.4
##   quality color
## 1      5   red
## 2      5   red
## 3      5   red
## 4      6   red
## 5      5   red
## 6      5   red
```

Data Summary

```
##   fixed.acidity  volatile.acidity  citric.acid  residual.sugar
## Min. : 3.800  Min. :0.0800  Min. :0.0000  Min. : 0.600
## 1st Qu.: 6.400 1st Qu.:0.2300 1st Qu.:0.2500 1st Qu.: 1.800
## Median : 7.000 Median :0.2900 Median :0.3100 Median : 3.000
## Mean   : 7.215 Mean   :0.3397 Mean   :0.3186 Mean   : 5.443
## 3rd Qu.: 7.700 3rd Qu.:0.4000 3rd Qu.:0.3900 3rd Qu.: 8.100
## Max.   :15.900 Max.   :1.5800 Max.   :1.6600 Max.   :65.800
##   chlorides  free.sulfur.dioxide total.sulfur.dioxide density
## Min. :0.00900 Min. : 1.00  Min. : 6.0  Min. :0.9871
## 1st Qu.:0.03800 1st Qu.:17.00 1st Qu.:77.0 1st Qu.:0.9923
## Median :0.04700 Median :29.00 Median :118.0 Median :0.9949
## Mean   :0.05603 Mean   :30.53 Mean   :115.7 Mean   :0.9947
## 3rd Qu.:0.06500 3rd Qu.:41.00 3rd Qu.:156.0 3rd Qu.:0.9970
## Max.   :0.61100 Max.   :289.00 Max.   :440.0 Max.   :1.0390
##   pH  sulphates  alcohol  quality
## Min. :2.720  Min. :0.2200  Min. : 8.00  Min. :3.000
## 1st Qu.:3.110 1st Qu.:0.4300 1st Qu.: 9.50 1st Qu.:5.000
## Median :3.210 Median :0.5100 Median :10.30 Median :6.000
## Mean   :3.219 Mean   :0.5313 Mean   :10.49 Mean   :5.818
```

```

## 3rd Qu.:3.320   3rd Qu.:0.6000   3rd Qu.:11.30   3rd Qu.:6.000
## Max.    :4.010   Max.    :2.0000   Max.    :14.90   Max.    :9.000
##      color
## Length:6497
## Class :character
## Mode  :character
##
##
##
```

** PCA using prcomp that uses the singular value decomposition (SVD)**

```

## Importance of components:
##                               PC1     PC2     PC3     PC4     PC5     PC6     PC7
## Standard deviation      1.7407 1.5792 1.2475 0.98517 0.84845 0.77930 0.72330
## Proportion of Variance 0.2754 0.2267 0.1415 0.08823 0.06544 0.05521 0.04756
## Cumulative Proportion  0.2754 0.5021 0.6436 0.73187 0.79732 0.85253 0.90009
##                               PC8     PC9     PC10    PC11
## Standard deviation      0.70817 0.58054 0.4772 0.18119
## Proportion of Variance 0.04559 0.03064 0.0207 0.00298
## Cumulative Proportion  0.94568 0.97632 0.9970 1.00000
```

** Interpreting the components **

```

##                               PC1     PC2
## fixed.acidity        -0.24  0.34
## volatile.acidity     -0.38  0.12
## citric.acid          0.15  0.18
## residual.sugar       0.35  0.33
## chlorides            -0.29  0.32
## free.sulfur.dioxide  0.43  0.07
## total.sulfur.dioxide 0.49  0.09
## density              -0.04  0.58
## pH                   -0.22 -0.16
## sulphates            -0.29  0.19
## alcohol              -0.11 -0.47
```

Summary of the loadings

```

##             Properties      PC1      PC2      PC3      PC4
## 1      fixed.acidity -0.23879890 0.33635454 -0.43430130 0.16434621
## 2      volatile.acidity -0.38075750 0.11754972 0.30725942 0.21278489
## 3      citric.acid    0.15238844 0.18329940 -0.59056967 -0.26430031
## 4      residual.sugar 0.34591993 0.32991418 0.16468843 0.16744301
## 5      chlorides       -0.29011259 0.31525799 0.01667910 -0.24474386
## 6      free.sulfur.dioxide 0.43091401 0.07193260 0.13422395 -0.35727894
## 7      total.sulfur.dioxide 0.48741806 0.08726628 0.10746230 -0.20842014
## 8      density          -0.04493664 0.58403734 0.17560555 0.07272496
## 9      pH                -0.21868644 -0.15586900 0.45532412 -0.41455110
## 10     sulphates         -0.29413517 0.19171577 -0.07004248 -0.64053571
## 11     alcohol           -0.10643712 -0.46505769 -0.26110053 -0.10680270
##             PC5      PC6      PC7      PC8      PC9      PC10
## 1   -0.1474804 -0.20455371 -0.28307944 0.401235645 0.3440567 -0.281267685
## 2    0.1514560 -0.49214307 -0.38915976 -0.087435088 -0.4969327 0.152176731
## 3   -0.1553487  0.22763380 -0.38128504 -0.293412336 -0.4026887 0.234463340
## 4   -0.3533619 -0.23347775  0.21797554 -0.524872935 0.1080032 -0.001372773
## 5    0.6143911  0.16097639 -0.04606816 -0.471516850 0.2964437 -0.196630217
```

```

## 6  0.2235323 -0.34005140 -0.29936325  0.207807585  0.3666563  0.480243340
## 7  0.1581336 -0.15127722 -0.13891032  0.128621319 -0.3206955 -0.713663486
## 8 -0.3065613  0.01874307 -0.04675897  0.004831136  0.1128800 -0.003908289
## 9 -0.4533764  0.29657890 -0.41890702 -0.028643277  0.1278367 -0.141310977
## 10 -0.1365769 -0.29692579  0.52534311  0.165818022 -0.2077642  0.045959499
## 11 -0.1888920 -0.51837780 -0.10410343 -0.399233887  0.2518903 -0.205053085
##          PC11
## 1 -0.3346792663
## 2 -0.0847718098
## 3  0.0011089514
## 4 -0.4497650778
## 5 -0.0434375867
## 6  0.0002125351
## 7  0.0626848131
## 8  0.7151620723
## 9 -0.2063605036
## 10 -0.0772024671
## 11  0.3357018784

##          Properties   PC1
## 1 total.sulfur.dioxide  0.49
## 2 free.sulfur.dioxide  0.43
## 3 residual.sugar      0.35
## 4 citric.acid         0.15
## 5 density              -0.04
## 6 alcohol              -0.11
## 7 pH                  -0.22
## 8 fixed.acidity        -0.24
## 9 chlorides            -0.29
## 10 sulphates           -0.29
## 11 volatile.acidity    -0.38

```

** PC1 seems to give more positive loadings to sulfur dioxide and residual sugar while chlorides, sulphates & acidity have negative loadings**


```

##          Properties   PC2
## 1          density  0.58
## 2          fixed.acidity  0.34
## 3          residual.sugar  0.33
## 4          chlorides  0.32
## 5          sulphates  0.19
## 6          citric.acid  0.18
## 7          volatile.acidity  0.12
## 8 total.sulfur.dioxide  0.09
## 9 free.sulfur.dioxide  0.07
## 10          pH -0.16
## 11          alcohol -0.47

```

PC2 gives more positive loading to density and high negative to alcohol content


** PC space, i.e. the space of summary variables we've created**



- We can see that there are two clusters emerging. One to the left having more negative PC1 & another to the right having more positive PC1
- As we do not have access to the wine color, we cannot say exactly which of these two clusters correspond to red & White wine
- However, the difference in the loadings for each feature by the principal components implies that these two clusters are distinct

Categorize quality

```
##  
## High Quality  Low Quality  
##          4113        2384
```



It doesn't look like the principal components can do a good job of separating high & low quality wines.

Clusters

```

##          fixed.acidity      volatile.acidity      citric.acid
##       6.85167903           0.27458385        0.33524928
##      residual.sugar      chlorides  free.sulfur.dioxide
##       6.39402555           0.04510424       35.52152864
## total.sulfur.dioxide      density            pH
##      138.45848785          0.99400486      3.18762464
##      sulphates            alcohol
##      0.48880511          10.52235888
##          fixed.acidity      volatile.acidity      citric.acid
##       8.2895922            0.5319416        0.2695435
##      residual.sugar      chlorides  free.sulfur.dioxide
##       2.6342666            0.0883238       15.7647596
## total.sulfur.dioxide      density            pH
##      48.6396835           0.9967404       3.3097200
##      sulphates            alcohol
##      0.6567194            10.4015216
##          volatile.acidity      chlorides      sulphates
##      -1.5631878           -1.2336600      -1.1284118
##      fixed.acidity          density            pH
##      -1.1091297           -0.9122427      -0.7593601
##      alcohol              citric.acid  residual.sugar

```

```

##          0.1013131      0.4521520      0.7902299
## free.sulfur.dioxide total.sulfur.dioxide
##          1.1130951      1.5890987

```

Table clusters


```

##
##      red white
## 1    24  4830
## 2   1575   68


## fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 1       7.4        0.70     0.00      1.9     0.076
## 2       7.8        0.88     0.00      2.6     0.098
## 3       7.8        0.76     0.04      2.3     0.092
## 4      11.2        0.28     0.56      1.9     0.075
## 5       7.4        0.70     0.00      1.9     0.076
## 6       7.4        0.66     0.00      1.8     0.075
## free.sulfur.dioxide total.sulfur.dioxide density pH sulphates alcohol
## 1           11        34  0.9978 3.51      0.56     9.4
## 2           25        67  0.9968 3.20      0.68     9.8
## 3           15        54  0.9970 3.26      0.65     9.8
## 4           17        60  0.9980 3.16      0.58     9.8
## 5           11        34  0.9978 3.51      0.56     9.4
## 6           13        40  0.9978 3.51      0.56     9.4
## quality color    dist1    dist2
## 1       5   red 107.40131 15.47603
## 2       5   red  72.34191 20.57183
## 3       5   red  87.02246  5.48984
## 4       6   red  80.86079 11.83767
## 5       5   red 107.40131 15.47603
## 6       5   red 101.11543  9.21505


```


Using the two most contrasting features for the 2 clusters as x & y to visualize





Plot for High and Low Quality Wine












** Summary**

- PCA is the best dimensionality reduction technique makes more sense to me for this data
- PCA Counteracts the issues of high-dimensional data
- PCA improves performance at a very low cost of model accuracy

Correlation Plot




The plot at first glance will just show us what how many and what variables are showing a correlation. We see that many of the variables are correlated. For example, the plot is showing us that college uni and online gaming are correlated. We also see that health nutrition and personal fitness are correlated as well.

We can use the principal component analysis to summarize the information content in large data tables by means of a smaller set of “summary indices” that can be more easily visualized and analyzed.

We've centered and scaled the data and extracted them which will be the attributes.


```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation 2.0662 1.63875 1.55225 1.47391 1.35444 1.2647 1.22840
## Proportion of Variance 0.1377 0.08663 0.07773 0.07008 0.05918 0.0516 0.04868
## Cumulative Proportion 0.1377 0.22435 0.30207 0.37215 0.43133 0.4829 0.53160
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation 1.18786 1.06740 0.99847 0.96967 0.95830 0.93741 0.91727
## Proportion of Variance 0.04552 0.03675 0.03216 0.03033 0.02962 0.02835 0.02714
## Cumulative Proportion 0.57712 0.61387 0.64603 0.67636 0.70599 0.73433 0.76147
##          PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation 0.89976 0.8571 0.8088 0.75303 0.69643 0.66873 0.65307
## Proportion of Variance 0.02611 0.0237 0.0211 0.01829 0.01565 0.01443 0.01376
## Cumulative Proportion 0.78759 0.8113 0.8324 0.85068 0.86632 0.88075 0.89451
##          PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation 0.64509 0.63596 0.61770 0.60923 0.59551 0.58874 0.55099
## Proportion of Variance 0.01342 0.01305 0.01231 0.01197 0.01144 0.01118 0.00979
## Cumulative Proportion 0.90793 0.92098 0.93329 0.94526 0.95670 0.96788 0.97767
##          PC29     PC30     PC31
## Standard deviation 0.51474 0.48583 0.43716
```

```
## Proportion of Variance 0.00855 0.00761 0.00616  
## Cumulative Proportion 0.98622 0.99384 1.00000
```



```
## [1] 0.6460315
```

Looking at the 10th principal component we can see that roughly 64.6% of the fraction of variance is explained. Using the Kaiser Criterion we can drop all Eigen values under 1.0.




Using the elbow method we've determine the number of clusters. We've decided to move forward with 10 clusters.

```

##      V1          V2          V3          V4
##  Min. :0.5477  Min. :0.3408  Min. :0.4142  Min. :0.4153
##  1st Qu.:0.8090 1st Qu.:0.5824 1st Qu.:0.5885 1st Qu.:0.6568
##  Median :1.0804 Median :0.8380 Median :0.7673 Median :0.8249
##  Mean   :1.3856 Mean  :1.5891 Mean  :0.9776 Mean  :1.6967
##  3rd Qu.:1.3894 3rd Qu.:1.2710 3rd Qu.:1.1911 3rd Qu.:1.3630
##  Max.   :9.2814 Max.  :11.0950 Max.  :3.3613 Max.  :11.1836
##      V5          V6          V7          V8
##  Min. :0.1760  Min. :0.4351  Min. :0.2504  Min. :0.2432
##  1st Qu.:0.3527 1st Qu.:0.7732 1st Qu.:0.5915 1st Qu.:0.5462
##  Median :0.4488 Median :0.9052 Median :0.7804 Median :0.8671
##  Mean   :0.5359 Mean  :1.5799 Mean  :1.1924 Mean  :1.3236
##  3rd Qu.:0.7067 3rd Qu.:1.3495 3rd Qu.:1.1384 3rd Qu.:1.1655
##  Max.   :1.2408 Max.  :11.6825 Max.  :6.8858 Max.  :6.7725
##      V9          V10
##  Min. :0.3933  Min. :0.4982
##  1st Qu.:0.7193 1st Qu.:0.6697
##  Median :1.0175 Median :0.9711
##  Mean   :1.5999 Mean  :1.4437
##  3rd Qu.:1.5344 3rd Qu.:1.5758
##  Max.   :6.1199 Max.  :7.1336


```

Cluster 1




Cluster 1 primarily composed with people who have interest in current_events, traveling and politics.

Cluster 2




Cluster 2 primarily composed with people who have interest in shopping, current_events, college_uni.

Cluster 3




Cluster 3 emphasized people who have interest in dating then followed by fashion.


Cluster 4




Cluster 4 primarily composed of people who have interest in sports_fandom, religion, and parenting.





Cluster 5 primarily composed of people who have interest in cooking, fashion and beauty.






Cluster 7 primarily composed of people who have interest in art, tv_film, and current_events.





Cluster 10



Cluster 10 primarily composed of people who have interest in personal_fitness, cooking, outdoors.

K-Means clustering can allow us to identify the different market segments. In this case, we used 10 Based on the K-Means clustering, we can identify distinct market segments that NutrientH2O can potentially leverage to design specific marketing campaigns.